

SW6100 Linux Wear OS Performance Technical Overview

80-74841-37 Rev. AA

Confidential – Qualcomm Technologies, Inc. and/or its affiliated companies – May Contain Trade Secrets

NO PUBLIC DISCLOSURE PERMITTED: Please report postings of this document on public servers or websites to: DocCtrlAgent@qualcomm.com.

© Qualcomm Technologies, Inc. and/or its subsidiaries. All rights reserved.

Revision History

Revision	Date	Description
AA	December 2025	Initial release

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Contents

1 Architecture	<u>5</u>
2 QTI Energy Awareness Scheduling Performance Features	<u>8</u>
3 Scheduler Update on SW6100	<u>10</u>
4 Scheduler Features on SW6100	<u>16</u>
5 Touch Boost	<u>29</u>
6 Governor (Walt)	<u>31</u>
7 Core Control (Halt)	<u>36</u>
8 Debugging	<u>42</u>

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Objectives

At the end of this presentation, you will understand the following performance-related aspects for this chipset:

- System architecture
- Scheduler new features and update.

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com



Section 1

Architecture

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyunchul.won@lge.com

SW6100 Applications Processor

- SW6100 has 1 × Arm® Cortex A78 Gold core:
 - 32 kB instruction cache, 32 kB data cache
 - 256 kB of L2 cache
 - Max frequency: 1.95 GHz
- SW6100 has 4 × Cortex A55 Silver cores:
 - 32 kB instruction cache, 32 kB data cache.
 - 64 kB of L2 cache
 - Max frequency: 2.1 GHz
- All cores share 3 MB of L3 cache.

CPU architecture

Gold	Cortex A78	
	32 KB/32KB1/D L1	
	256 KB L2	
Silver	Cortex A55	Cortex A55
	32 KB/32 KB I/D L1	32 KB/32 KB I/D L1
	64 KB L2	64 KB L2
	Cortex A55	Cortex A55
	32 KB/32 KB I/D L1	32 KB/32 KB I/D L1
	64 KB L2	64 KB L2

3 MB L3 Cache

Application Processor Performance Overview

Features	Description
Process	3nm process node
CPU SS	1 × A78 2.1 GHz (Arm Cortex), 256 kB L2 cache, 32 kB L1 cache 4 × A55 1.95 GHz (Arm Cortex), 64 kB L2 cache, 32 kB of dedicated L1 cache 3 MB L3 cache
DDR SS	DDR with 128 KB LLC, LP5 @ 3.2 GHz and LP5.x @ 4.2 GHz support
GPU SS	Qualcomm® Adreno™ 622 v1 graphics

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com



Section 2

QTI Energy Awareness Scheduling Performance Features

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyunchul.won@lge.com

QTI Energy Awareness Scheduling

- QTI energy awareness scheduling (EAS) is designed to fulfill the market requirements.
- Key changes in the QTI energy awareness scheduling (EAS):
 - It is upgraded to Kernel 6.12
 - Walt support is enabled for 2* cluster CPU architecture on SW6100
 - Silver cluster (APC0)
 - Gold cluster (APC1)

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com



Section 3

Scheduler Update on SW6100

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

WALT Highlight on SW6100

- Task migration:
 - Per cgroup up/down migrate
- Cpufreq governor:
 - Zone based target load

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Task Migration: Per cgroup sched_up/down Migration

- New cgroup-based force migration is introduced, which provides better granularity of task migration control.
- Cgroup support: topapp, foreground, background, and other; and it is per cluster setting.
 - `/proc/sys/walt/clusterX/sched_topapp_updownmigrate`
 - `/proc/sys/walt/clusterX/sched_foreground_updownmigrate`
 - `/proc/sys/walt/clusterX/sched_background_updownmigrate`
 - `/proc/sys/walt/clusterX/sched_other_cgroup_updownmigrate`
- {upmigrate,downmigrate} works as a pair. The first entry is the upmigrate percentage and the second entry is the downmigrate percentage.
- Example:
 - `echo 95 85 > /proc/sys/walt/cluster0/sched_background_updownmigrate`
 - 95: upmigrate threshold; 85: downmigrate threshold. Both are in percentage.
- Perflock is supported
- Note:
 - `sched_other_cgroup_updownmigrate` will be the “default” one. That is, if one task is not in topapp/foreground/background cgroup. It will be controlled by `sched_older_cgroup_updownmigrate`.
 - Legacy upmigrate/downmigrate are outdated and deprecated..
 - `/proc/sys/walt/sched_downmigrate/sched_upmigrate`
 - `/proc/sys/walt/sched_early_downmigrate/sched_early_upmigrate`

Sched Upmigrate/Downmigrate tunable

- Provides a threshold to migrate a task to the right cluster to improve UX performance.
- From SW6100, we provide cgroup based task force migration as below table.
- Sched group upmigrate/downmigrate is disabled by default.
- Lower value is performance oriented, and the higher value is power-oriented.

Name	Default value	Unit	Perflock opcode
sched_topapp_updownmigrate	95 85	Percentage	MPCTLV3_SCHED_TA_UPDOWNMIGRATE (0x43C80000)
sched_foreground_updownmigrate	95 85	Percentage	MPCTLV3_SCHED_FG_UPDOWNMIGRATE (0x43C78000)
sched_background_updownmigrate	95 85	Percentage	MPCTLV3_SCHED_BG_UPDOWNMIGRATE (0x43C74000)
sched_other_cgroup_updownmigrate	95 85	Percentage	MPCTLV3_SCHED_OT_UPDOWNMIGRATE (0x43C7C000)
sched_group_upmigrate	0		MPCTLV3_SCHED_GROUP_UPMIGRATE (0x40C54000)
sched_group_downmigrate	0		MPCTLV3_SCHED_GROUP_DOWNMIGRATE (0x40C58000)
sched_coloc_downmigrate_ns	400000000		MPCTLV3_SCHED_COLOC_DOWNMIGRATE_NS (0x40C8C000)

NOTE:

- This section will be updated later with the latest clock plan.
- Early up/down migration is deprecated

Waltgov: Zone-based Target Load

- This feature specifies the frequency of zone and the equivalent target load percentage to jump to next frequency.
- Configurable as {freq, percentage} zones.
- Example:
 - `echo 883200 95 1190400 90 > /sys/devices/system/cpu/cpufreq/policy0/walt/zone_max_util_pct`
 - above means next frequency will be range of
 - 1 ~ 883M if load is reaching 95%;
 - 884M ~ 1.19 G if load is reaching 90%
 - Other 1.19 G ~ Fmax if load is reaching 80%
- Perflock supported
- **Note:**
 - Max to 10 zones supported.
 - Target load percentage is range of {80,100}
 - This should be in descending order for each zone.
 - For example: illegal setting: 1190400 95 883200 90
 - Zones and target load (percentage) can change during runtime.
 - `/sys/devices/system/cpu/cpufreq/policy0/walt/zone_max_util_pct`
 - If the user updates the number of zones and then sets the corresponding target load values, any previously configured values for the old zones will be reset.
 - `target_load_thresh/target_load_shift` are outdated and deprecated.
 - Perflock support: `MPCTLV3_ZONE_MAX_UTIL_PCT(0x6081C000)`
 - For example: `0x6081C000, 4, 883200, 95, 1190400, 90`

Tunable

Name	zone_max_util_pct
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/zone_max_util_pct
Default value	2147483647 80
Description	{freq, percentage} zones
Perflock	MPCTLV3_ZONE_MAX_UTIL_PCT(0x6081C000)

- Configurable as {freq, percentage} zones.
- Example : `echo 883200 95 1190400 90 > /sys/devices/system/cpu/cpuX/cpufreq/walt/zone_max_util_pct`
- Perflock support: `MPCTLV3_ZONE_MAX_UTIL_PCT(0x6081C000)`
 - For example: `0x6081C000,4, 883200,95, 1190400,90`



Section 4

Scheduler Features on SW6100

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyunchul.won@lge.com

Sched Boost (CGROUP Boost)

- Aids groups of tasks to run on performance-friendly cores.
- Variants:
 - Full Throttle boost: All tasks promoted to all CPUs
 - Conservative boost – No force upmigrate is required, only enable early detection to address the latency sensitive case.
 - Restrained boost – No task promotions are made in this variant; distributed load from top app tasks drives CPU frequency.
 - Storage boost – Provides more CPU resources to meet the high I/O throughput case.
- Perflock opcode – `MPCTLV3_SCHED_BOOST(0x40C00000)`
- Sysfs node: `/proc/sys/walt/sched_boost`

Boost type	Boost value	Colocation	Frequency aggregate	Force upmigrate	Early detection	C-state and power cost considerations	Core control*
Full throttle	1	Top app	Yes	All tasks	Yes	No	max_cpus*
Conservative	2	No	No	No	Yes	■ No – Top app and foreground tasks ■ Yes – Other tasks	Default
Restrained	3	Top app	Top app	No	No	Yes	Default
Storage boost	4	No	No	No	No	No	max_cpus**
Balance boost	5	Top app	No	No	Yes	Yes	max_cpus*
Off	0	Top app	No	No	No	Yes	Default

NOTE:

* Core CTL configuration may be changed accordingly.

** If max_cpus is limited by some reason, like thermal, sched_boost(1) will not be able to use all the CPUs.

Early Detect

- Early detection helps to reduce the latency to report the load for some latency-sensitive cases.
 - Only for boost case (boost policy enable)
 - Task wake-up time gap away from last wake-up is longer than 9.5 ms.
- New tunable

Tunable	Comments	Perflock opcode
sched_ed_boost	■ /proc/sys/walt/sched_ed_boost ■ Percentage; if the default is 10 then it gives 110% boost to load	–

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Storage Boost

- If there is a higher QoS request from a UFS driver, a storage boost is enabled to provide more CPU resources to get better I/O performance.
- Storage boost enables boost for core CTL and unhalt all the possible CPUs.
- During storage boost, the fmax cap is not applied.

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Task-oriented Scheduling

- Apart from the CGROUP-based boost, a smaller granularity schedule policy is also provided to get more precise scheduling accuracy with the following tunable:

`/proc/sys/walt/sched_task_read_pid` gives the scheduler an ability to read the task ID*

Tunable	Comments	Perflock opcode
<code>sched_per_task_boost</code>	<code>sched_per_task_boost</code> is to boost the specific task to gold CPU using the PID ■ Boost to Mid – Disabled and does not work on SW6100 ■ Boost to Max – Task should run on Gold CPU ■ Strict Boost to Max – Task should run on Gold, even if Gold gets overutilized.	MPCTLV3_SCHED_TASK_BOOST (0x40C80000)
<code>sched_per_task_boost_period_ms</code>	Boost duration	–
<code>sched_init_task_load</code>	Give task init load when the task is forked.	MPCTLV3_SCHED_INIT_TASK_LOAD (0x40C18000)
<code>sched_low_latency</code>	Give latency sensitive CFS task a high possibility to be picked to run	MPCTLV3_SCHED_LOW_LATENCY (0x40CD0000)
<code>sched_wake_up_idle</code>	Wake up task running on an idle CPU	MPCTLV3_SCHED_WAKE_UP_IDLE (0x40CF0000)
<code>task_load_boost*</code>	To give a specific task load boost (rang -90%~90%)	This information will be updated in a later revision of this document.
<code>task_reduce_affinity</code>	To give a specific RT task like SurfaceFlinger to reduce the affinity to avoid power impact. It is not needed on SW6100.	–

NOTE:

- Usage sequence – Read the task ID and then set the per task tunable. Ensure that you disable the setting explicitly or use perflock, which passes the PID.
- *As group migration is disabled by default, the colocated task features are also disabled by default.

LPM Bias

- Sched-driven Idle mode guidance.
- Big cores usually have large private caches; caches must be flushed when the core enters deep Low-power modes.
- It takes hundreds of microseconds to bring BIG core out of the power collapse.
- Performance bottleneck occurs since BIG cores take a long time to respond and load. The power bottleneck (increased cache flushing activity) occurs if BIG cores are constantly entering/exiting deep LPMs.
- Sched code tracks various parameters on the CPU and guides the idle CPU governor to hold CPUs in shallow Low-power modes.

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Busy Hysteresis Tunables

Tunable	Comment	Unit	Perflock
sched_busy_hysteresis_enable_cpus	CPU mask contributes to the busy hysteresis	CPU mask	MPCTLV3_SCHED_BUSY_HYSTERESIS_CPU_MASK
sched_busy_hyst_ns	Default CPU busy hysteresis time	ns	MPCTLV3_SCHED_COLOC_BIAS_HYST
sched_coloc_busy_hysteresis_enable_cpus*	Colocate CPU mask contributes to the busy hysteresis	Colocate CPU mask	MPCTLV3_SCHED_BUSY_HYSTERESIS_ENABLE_COLOC_CPUS
sched_coloc_busy_hyst_max_ms*	Maximum colocate time; if the RTB time is longer than the maximum, update the hysteresis directly.	ms	—
sched_coloc_busy_hyst_cpu_busy_pct*	Percentage of per CPU cap	Percentage	MPCTLV3_SCHED_COLOC_BUZY_HYST_CPU_PCT
sched_coloc_busy_hyst_cpu_ns*	Per CPU hysteresis busy time	ns	MPCTLV3_SCHED_COLOC_BUZY_HYST_CPU_NS
sched_util_busy_hyst_cpu_util	Util of per CPU	—	MPCTLV3_SCHED_UTIL_BUSY_HYST_CPU_NS_BIG

NOTE: * As group migration is disabled by default, the colocated task features are disabled by default.

Task Filter

Name	sched_min_task_util_for_boost	Name	sched_min_task_util_for_colocation
Location	/proc/sys/walt/sched_min_task_util_for_boost	Location	/proc/sys/walt/sched_min_task_util_for_colocation
Default value	51	Default value	35
Description	Unit in percentage	Description	Unit in percentage
Perflock opcode	MPCTLV3_SCHED_MIN_TASK_UTIL_FOR_BOOST(0x40C7C000)	Perflock opcode	MPCTLV3_SCHED_MIN_TASK_UTIL_FOR_COLOCATION(0x40C78000)

- These sched_min_task_util_for_boost and sched_min_task_util_for_colocation tunables determine if:
 - A task is in colocation
 - The task with demand is **lower** than the /proc/sys/walt/sched_min_task_util_for_colocation threshold, then it stays on the Silver cluster.
 - The task with demand is **higher** than the /proc/sys/walt/sched_min_task_util_for_boost threshold, it would be a boost on to the Gold clusters.

Name	sched_task_unfilter_period
Location	/proc/sys/walt/sched_task_unfilter_period
Default value	100000000
Description	Unit in ns
Perflock opcode	MPCTLV3_SCHED_TASK_UNFILTER_PERIOD(0x40CB8000)

- The sched_task_unfilter_period tunable monitors the task filter valid time. The default time is 100 ms.
- Lower value is performance-oriented and the higher value is power-oriented.
- sched_min_task_util_for_uclamp for migration is disabled on SW6100.

sched_many_wakeup_threshold

- In a few use cases, a task wakes the other tasks back and forth, which crowds a single cluster. The `sched_many_wakeup_threshold` tunable provides the threshold to avoid such scenarios.
- If the wake-up task count on a specific CPU is greater than the threshold, the next CPU is skipped.
- The default value is 1000.
- **Perflock opcode:** `MPCTLV3_SCHED_MANY_WAKEUP_THRESHOLD` (0x40C9C000)

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com


sched_user_hint

- The change attempts to boost the top-app hosting cluster frequency.
- The sysnode gives the top-app load to trigger the frequency boost:

$$load = \frac{load * sched_user_hint}{100}$$

- This hint rests to 0 after 1 sec.
- Hint range is 0 to 1000.
- Perflock opcode: `MPCTLV3_SCHED_USER_HINT(0x40C84000)`

Task Packing Control



- Packing tasks to the same CPU benefits power but may impact performance. There must be a balance between performance and power for some specific cases.
- Here are some conditions that can control the packing fast path:
 - Packing tasks in the same cluster, if possible.
 - Packing happens only on unhalted CPUs.
 - Packing is not possible:
 - If `sched_idle_enough` = 0 (false), the fast path packing is disabled
 - If the current cluster load is over `sched_cluster_util_thres_pct`
 - High load CPU (higher than `sched_idle_enough`) does not accommodate packing
 - High demand task (higher than `sched_idle_enough`) is packing
 - When the average load is dropping
 - Packing is not possible on partial halt CPU
 - For any other use cases, packing tasks can happen to save power:

Tunable	Comment	Unit	
<code>sched_idle_enough</code>	Threshold to decide possible packing: ■ High demand task ■ High util CPU	Task demand	<code>MPCTLV3_SCHED_IDLE_ENOUGH</code>
<code>sched_idle_enough_clust</code>	Threshold to decide possible packing: ■ High util cluster ■ High util CPU	CPU util	<code>MPCTLV3_SCHED_IDLE_ENOUGH_CLUST_BIG</code> <code>MPCTLV3_SCHED_IDLE_ENOUGH_CLUST_LITTLE</code>
<code>sched_cluster_util_thres_pct</code>	Threshold to decide cluster is high load or not	Percentage	<code>MPCTLV3_SCHED_CLUSTER_UTIL_THRES_PCT</code>
<code>sched_cluster_util_thres_pct_clust</code>	Threshold (per cluster) to decide which cluster is high load or not	—	<code>MPCTLV3_SCHED_CLUSTER_UTIL_THRES_PCT</code>

Smart Frequency Management

- Workload with low IPC will see less performance improvement from using higher frequencies but can still burn high power.
- Frequency ceiling management
 - A smart frequency algorithm considers both the IPC and other workload characteristics to determine the appropriate frequency ceiling for the CPU.



- Algorithm re-evaluates the workload every sched tick (4 ms) to determine the correct frequency ceiling level on each cluster.
 - For example, in a cluster with two CPU cores, the 1st core ceiling is fixed to 1.8 GHz, and the second core ceiling is fixed to 2 GHz, then max ceiling is chosen for that cluster.
- Workload based frequency ceiling selection input
 - Frequency ceiling determination also considers other aspects beyond IPC, such as # misfits in the system, wake up count, as well as sustained high-load conditions.
 - Higher frequency can still be required even with lower IPC for scenarios with heavy load and low-performance tolerance. In these cases, the IPC-based frequency ceiling can be overwritten.

Smart Frequency Management - Tunable

- **IPC based frequency ceiling selection input**
 - IPC zones are defined to represent workload efficiency, and each IPC zone has a corresponding frequency ceiling.
 - The highest IPC among the CPUs in cluster will set the frequency ceiling based on preset thresholds.
- **Ways to uncap the IPC based frequency ceiling:**
 - use `scaling_min_freq` to boost
 - use full throttle boost
 - update `ipc_reason_config[level].freq_allowed`
- **Debug:**
 - `trace_ipc_update`
 - `trace_ipc_freq`
 - `trace_sched_freq_uncap`
 - `trace_waltgov_util_update`
 - `trace_waltgov_next_freq`

Tunable	Default value	Unit/Parameters	Perflock
/proc/sys/walt/clusterX/smart_freq/ipc_freq_levels	For example: 2400000 2400000 2700000 3000000 2147483647	Frequency	MPCTLV3_SCHED_IPC_LEVELS_BIG(0x63C50000)
/proc/sys/walt/clusterX/smart_freq/legacy_freq_levels	0:10:2147483647:0:0	{cluster,Index, frequency, Hysteresis , enable/Disable}	MPCTLV3_SCHED_LEGACY_FREQ_LEVELS(0x63C5C000)
/proc/sys/walt/cluster*/smart_freq/sched_smart_freq_dump_ipc_reason	0:0:0:2147483647:0:0 0:1:0:2147483647:0:0 0:2:0:2147483647:0:0 0:3:0:2147483647:0:0 0:4:0:2147483647:0:0 1:0:0:2147483647:0:0	{cluster,Index,IPC Value, frequency, Hysteresis , enable/Disable}	NA



Section 5

Touch Boost

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

About Touch Boost

- Touch Boost feature gives a system boost when the touch event is received. This provides better UX performance for the improved input latency.
- To follow the GKI change, the Touch Boost-related nodes should be moved to the same folder as Walt.

Path	Kernel 5.4	Kernel 5.10 and later versions
Code	techpack/sched/cpu-boost.c	kernel/sched/walt/input-boost.c
Tunable	/sys/devices/system/cpu/cpu_boost/	/proc/sys/walt/input_boost/

- Tunable:

Tunable	Comments	Format
input_boost_freq	Boost frequency	{cpu:freq} Example: 0:960000 1:0 2:0 3:0 4:0 5:0
input_boost_ms	Boost duration	In ms

- Lower (short) value is power-oriented and the higher (longer) value is performance-oriented.
- Perflock opcode `MPCTLV3_INPUT_BOOST_RESET(0x42C08000)` support is to change the input boost frequency for some specific cases.



Section 6

Governor (Walt)

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

About Walt Governor

- SchedUtil Governor is no longer used after Kernel 5.10. Instead, Walt Governor is used for the CPU frequency guidance.
- Walt works with the EAS to change the CPU frequency. It also supports `fast_switch` by default.
- Cpufreq driver follows GKI upstream SCMI cpufreq driver, which is specification-compliant.
- `fast_switch` does the following:
 - Facilitates frequency switching from the scheduler's context without spawning the `waltgov` thread.
 - Reduces the latency of the switching frequency.


Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Walt Tunable – adaptive_freq (CPU Load Aware Boost)

- Added adaptive_level_1 to minimize possible power impact from loss of Silver cores during lightly-loaded portion of Launch
- These tunables give frequency bucket in some use cases:
 - If the vote frequency < adaptive_level_1, target frequency is adaptive_level_1.
 - If the vote frequency is between adaptive_level_1 and adaptive_low_freq, the target frequency is adaptive_low_freq.
 - If the vote frequency is between adaptive_low_freq and adaptive_high_freq, the target frequency is adaptive_high_freq.
- This avoids governor voting latency when the load is not high enough.
- Lower value is power-oriented and higher value is performance-oriented.

Name	adaptive_high_freq
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/adaptive_high_freq
Default value	0 (This information will be updated in a future revision of this document)
Description	Unit in Hz
Perflock code	MPCTLV3_WALT_ADAPTIVE_HIGH_FREQ_BIG(0x4080C000)

Name	adaptive_level_1
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/adaptive_level_1
Default value	0 (This information will be updated in a future revision of this document)
Description	Unit in Hz
Perflock code	MPCTLV3_WALT_ADAPTIVE_LEVEL_1_FREQ_BIG(0x40814000)




Name	adaptive_low_freq
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/adaptive_low_freq
Default value	0 (This information will be updated in a future revision of this document)
Description	Unit in Hz
Perflock code	MPCTLV3_WALT_ADAPTIVE_LOW_FREQ_BIG(0x40808000)

Walt Tunable - High-speed

Name	hispeed_freq
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/hispeed_freq
Default value	1036800/1440000
Description	- Frequency bump to reach hispeed_load - Lower value is power-oriented, higher value is performance-oriented
Perflock	MPCTLV3_SCHEDUTIL_HISPEED_FREQ_CLUSTER_BIG (0x4143C000)

Name	hispeed_cond_freq
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/hispeed_cond_freq
Default value	0
Description	- Frequency threshold - If avg_load is dropping below the threshold, hiloal can be break
Perflock	MPCTLV3_WALT_HISPEED_COND_FREQ_BIG (0x40818000)

Name	hispeed_load
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/hispeed_load
Default value	90/90
Description	- Percentage load to bump to hispeed_freq - Lower value is power-oriented, higher value is performance-oriented
Perflock	MPCTLV3_SCHEDUTIL_HISPEED_LOAD_CLUSTER_BIG (0x41440000)



- If the average load reaches hispeed_load, the frequency bumps up to hispeed_freq
- Reduces the latency of CPU frequency ramping.
- Tunable pairs match the cluster numbers.
- [New] hispeed_cond_freq will give the change that break the hiloal condition if the avg_load is dropping
- Lower value is power-oriented and higher value is performance-oriented.
Load will be avg – top app load

NOTE: * This section would be updated later with clock plan details.

Walt Tunable – rtg_boost

Name	rtg_boost_freq
Location	/sys/devices/system/cpu/cpuX/cpufreq/walt/rtg_boost_freq
Default value	249600/960000
Description	Frequency
Perflock opcode	MPCTLV3_SCHEDUTIL_RTG_BOOST_FREQ_CLUSTER_BIG(0x4144C000)

- This tunable can give higher boost input for specific clusters when the target CPU frequency for the colocated group tasks is selected.
- When colocation is active, that is, top application is preferred and the cluster is Gold, the CPU frequency is pegged at `rtg_boost_frequency` minimum. At present, it is used only for the Silver cluster so that filtered tasks from the Gold cluster can use the higher frequency.
- Lower value is power-oriented and higher value is performance-oriented.

NOTE: Ensure that you enable the group upmigrate at the same time when you set `rtg_boost`. * This section would be updated later with the clock plan details.



Section 7

Core Control (Halt)

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyunchul.won@lge.com

Introduction

- No partial halt on SW6100.
- **Node:** /sys/devices/system/cpu/cpuX/core_ctl/

Parameter	Silver	Gold
core_ctl enable	1	1
min_cpus	1	0
Max_cpus	4	1
not_preferred	0 1 1 1	0

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Halt Tunable

Name	busy_down_thres
Location	/sys/devices/system/cpu/cpuX/core_ctl/busy_down_thres
Default value	30 30 30 30 (Silver)/30 (Gold)
Description	Unit in percentage

Name	busy_up_thres
Location	/sys/devices/system/cpu/cpuX/core_ctl/busy_up_thres
Default value	60 60 60 60 (Silver)/60 (Gold)
Description	Unit in percentage

- If the load of the CPU is under the busy down threshold, the CPU is marked as not busy.
- If the load is above the busy up threshold, the CPU is marked as busy and may need more CPU online.
- Lower value is performance-oriented and higher value is power-oriented.

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@qcom.com

Halt Tunable (cont.)

Name	min_cpus
Location	/sys/devices/system/cpu/cpuX/core_ctl/min_cpus
Default value	1 (Silver)/ 0 (Gold)
Description	Core number
Perflock	▪ MPCTLV3_MIN_ONLINE_CPU_CLUSTER_BIG(0x41000000)

Name	max_cpus
Location	/sys/devices/system/cpu/cpuX/core_ctl/max_cpus
Default value	4 (Silver)/1 (Gold)
Description	Core number
Perflock	▪ MPCTLV3_MAX_ONLINE_CPU_CLUSTER_BIG(0x41004000)

Halt Tunable (cont.)

Name	Task_thres
Location	/sys/devices/system/cpu/cpuX/core_ctl/task_thres
Default value	4(Silver)/1(Gold)
Description	Heavy task number

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

NOTE: If the tasks on the cluster are more than the task_thres, all the cores on the cluster are unisolated.

offline_delay_ms

Name	offline_delay_ms
Location	/sys/devices/system/cpu/cpuX/core_ctl/offline_delay_ms
Default value	1000/1000 (ms)
Description	This tunable allows to change the CPU halt/unhalt behavior.

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com



Section 8

Debugging

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyunchul.won@lge.com

EAS Trace Points in Kernel 6.12

- All the upstream sched traces are enabled using "sched:*" name

```
adb shell "echo sched:* >> /sys/kernel/tracing/set_event"
```

- Example:

- adb shell "echo sched:sched_waking >> /sys/kernel/tracing/set_event"
- adb shell "echo sched:sched_wakeup >> /sys/kernel/tracing/set_event"

- Walt traces are enabled by using "schedwalt:*" name

```
adb shell "echo schedwalt:* >> /sys/kernel/tracing/set_event"
```

- Example:

- adb shell "echo schedwalt:sched_find_best_target >> /sys/kernel/tracing/set_event"
- adb shell "echo schedwalt:sched_task_util >> /sys/kernel/tracing/set_event"

Acronyms and Terms

Acronym or term	Definition
DCVS	Dynamic clock and voltage scaling
EAS	Energy aware scheduling

Qualcomm
Confidential - May Contain Trade Secrets
2026-05-14 09:30:33 GMT
hyungchul.won@lge.com

Questions?

For additional information or to submit technical questions, go to: <https://www.qualcomm.com/support>

LEGAL INFORMATION

Your access to and use of this material, along with any documents, software, specifications, reference board files, drawings, diagnostics and other information contained herein (collectively this “Material”), is subject to your (including the corporation or other legal entity you represent, collectively “You” or “Your”) acceptance of the terms and conditions (“Terms of Use”) set forth below. If You do not agree to these Terms of Use, you may not use this Material and shall immediately destroy any copy thereof.

1) Legal Notice.

This Material is being made available to You solely for Your internal use with those products and service offerings of Qualcomm Technologies, Inc. (“Qualcomm Technologies”), its affiliates and/or licensors described in this Material, and shall not be used for any other purposes. If this Material is marked as “Qualcomm Internal Use Only”, no license is granted to You herein, and You must immediately (a) destroy or return this Material to Qualcomm Technologies, and (b) report Your receipt of this Material to qualcomm.support@qti.qualcomm.com. This Material may not be altered, edited, or modified in any way without Qualcomm Technologies’ prior written approval, nor may it be used for any machine learning or artificial intelligence development purpose which results, whether directly or indirectly, in the creation or development of an automated device, program, tool, algorithm, process, methodology, product and/or other output. Unauthorized use or disclosure of this Material or the information contained herein is strictly prohibited, and You agree to indemnify Qualcomm Technologies, its affiliates and licensors for any damages or losses suffered by Qualcomm Technologies, its affiliates and/or licensors for any such unauthorized uses or disclosures of this Material, in whole or part.

Qualcomm Technologies, its affiliates and/or licensors retain all rights and ownership in and to this Material. No license to any trademark, patent, copyright, mask work protection right or any other intellectual property right is either granted or implied by this Material or any information disclosed herein, including, but not limited to, any license to make, use, import or sell any product, service or technology offering embodying any of the information in this Material.

THIS MATERIAL IS BEING PROVIDED “AS IS” WITHOUT WARRANTY OF ANY KIND, WHETHER EXPRESSED, IMPLIED, STATUTORY OR OTHERWISE. TO THE MAXIMUM EXTENT PERMITTED BY LAW, QUALCOMM TECHNOLOGIES, ITS AFFILIATES AND/OR LICENSORS SPECIFICALLY DISCLAIM ALL WARRANTIES OF TITLE, MERCHANTABILITY, NON-INFRINGEMENT, FITNESS FOR A PARTICULAR PURPOSE, SATISFACTORY QUALITY, COMPLETENESS OR ACCURACY, AND ALL WARRANTIES ARISING OUT OF TRADE USAGE OR OUT OF A COURSE OF DEALING OR COURSE OF PERFORMANCE. MOREOVER, NEITHER QUALCOMM TECHNOLOGIES, NOR ANY OF ITS AFFILIATES AND/OR LICENSORS, SHALL BE LIABLE TO YOU OR ANY THIRD PARTY FOR ANY EXPENSES, LOSSES, USE, OR ACTIONS HOWSOEVER INCURRED OR UNDERTAKEN BY YOU IN RELIANCE ON THIS MATERIAL.

Certain product kits, tools and other items referenced in this Material may require You to accept additional terms and conditions before accessing or using those items.

Technical data specified in this Material may be subject to U.S. and other applicable export control laws. Transmission contrary to U.S. and any other applicable law is strictly prohibited.

Nothing in this Material is an offer to sell any of the components or devices referenced herein.

This Material is subject to change without further notification.

In the event of a conflict between these Terms of Use and the *Website Terms of Use* on www.qualcomm.com, the *Qualcomm Privacy Policy* referenced on www.qualcomm.com, or other legal statements or notices found on prior pages of the Material, these Terms of Use will control. In the event of a conflict between these Terms of Use and any other agreement (written or click-through, including, without limitation any non-disclosure agreement) executed by You and Qualcomm Technologies or a Qualcomm Technologies affiliate and/or licensor with respect to Your access to and use of this Material, the other agreement will control.

These Terms of Use shall be governed by and construed and enforced in accordance with the laws of the State of California, excluding the U.N. Convention on International Sale of Goods, without regard to conflict of laws principles. Any dispute, claim or controversy arising out of or relating to these Terms of Use, or the breach or validity hereof, shall be adjudicated only by a court of competent jurisdiction in the county of San Diego, State of California, and You hereby consent to the personal jurisdiction of such courts for that purpose.

2) Trademark and Product Attribution Statements.

Qualcomm is a trademark or registered trademark of Qualcomm Incorporated. Arm is a registered trademark of Arm Limited (or its subsidiaries) in the U.S. and/or elsewhere. The Bluetooth® word mark is a registered trademark owned by Bluetooth SIG, Inc. Other product and brand names referenced in this Material may be trademarks or registered trademarks of their respective owners.

Snapdragon and Qualcomm branded products referenced in this Material are products of Qualcomm Technologies, Inc. and/or its subsidiaries. Qualcomm patented technologies are licensed by Qualcomm Incorporated.

THE DOCUMENTATION ACCOMPANYING THE MATERIALS AND/OR RELEVANT PRODUCTS AND SERVICE OFFERINGS MAY INCLUDE IMPORTANT USE LIMITATIONS. ANY DEVIATIONS FROM APPLICABLE USE LIMITATIONS MAY ADVERSELY IMPACT PERFORMANCE, DURABILITY, QUALITY OR SAFETY. YOU ASSUME ALL RISKS AND LIABILITIES ASSOCIATED WITH ANY DEVIATIONS.